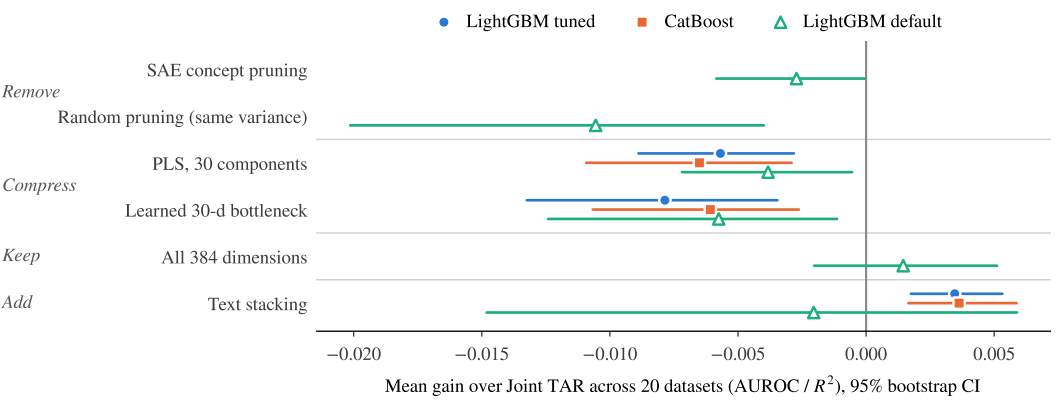
LightGBM tuned
CatBoost
LightGBM default
SAE concept pruning
Remove
Random pruning (same variance)
PLS, 30 components
Compress
Learned 30-d bottleneck
Keep
All 384 dimensions
Add
Text stacking
−0.020
−0.015
−0.010
−0.005
0.000
0.005
Mean gain over Joint TAR across 20 datasets (AUROC / R^2), 95% bootstrap CI